

Q. Write a program to implement Boosting and Bagging methods. Analyze using suitable dataset.

Importing necessary library

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import BaggingClassifier, AdaBoostClassifier,
GradientBoostingClassifier, RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix, precision_score, recall_score, f1_score
from sklearn.preprocessing import LabelEncoder
```

Load the dataset and displaying head of the dataset.

```
df = pd.read_csv('Almond.csv')
df.head()
```

	Unnamed: 0	Length (major axis)	Width (minor axis)	Thickness (depth)	\
0	0	NaN	227.940628	127.759132	
1	1	NaN	234.188126	128.199509	
2	2	NaN	229.418610	125.796547	
3	3	NaN	232.763153	125.918808	
4	4	NaN	230.150742	107.253448	

	Area	Perimeter	Roundness	Solidity	Compactness	Aspect Ratio	\
0	22619.0	643.813269	NaN	0.973384	1.458265	NaN	
1	23038.0	680.984841	NaN	0.957304	1.601844	NaN	
2	22386.5	646.943212	NaN	0.967270	1.487772	NaN	
3	22578.5	661.227483	NaN	0.965512	1.540979	NaN	
4	19068.0	624.842706	NaN	0.951450	1.629395	NaN	

	Eccentricity	Extent	Convex hull (convex area)	Type
0	NaN	0.681193	23237.5	MAMRA
1	NaN	0.656353	24065.5	MAMRA
2	NaN	0.683620	23144.0	MAMRA
3	NaN	0.685360	23385.0	MAMRA
4	NaN	0.714800	20041.0	MAMRA

Displaying info

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2803 entries, 0 to 2802
```

```
Data columns (total 14 columns):
```

```
#    Column                                Non-Null Count  Dtype
```

```

-----
0    Unnamed: 0                2803 non-null    int64
1    Length (major axis)      1946 non-null    float64
2    Width (minor axis)       1861 non-null    float64
3    Thickness (depth)        1799 non-null    float64
4    Area                     2803 non-null    float64
5    Perimeter                2803 non-null    float64
6    Roundness                1946 non-null    float64
7    Solidity                 2803 non-null    float64
8    Compactness              2803 non-null    float64
9    Aspect Ratio             1004 non-null    float64
10   Eccentricity             1004 non-null    float64
11   Extent                   2803 non-null    float64
12   Convex hull(convex area) 2803 non-null    float64
13   Type                     2803 non-null    object
dtypes: float64(12), int64(1), object(1)
memory usage: 306.7+ KB

```

Checking for null values

```

df.isnull().sum()

Unnamed: 0                0
Length (major axis)       857
Width (minor axis)        942
Thickness (depth)         1004
Area                      0
Perimeter                 0
Roundness                 857
Solidity                  0
Compactness               0
Aspect Ratio              1799
Eccentricity              1799
Extent                    0
Convex hull(convex area)  0
Type                     0
dtype: int64

```

Plotting Null values

```

# Plotting missing values heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Values Heatmap')
plt.show()


# Summary statistics
summary_stats = df.describe()

# Distribution of the 'Type' column
type_distribution = df['Type'].value_counts()

summary_stats, type_distribution

```

Missing Values Heatmap



	Unnamed: 0	Length (major axis)	Width (minor axis)	\
count	2803.000000	1946.000000	1861.000000	
mean	1401.000000	290.609274	171.025915	
std	809.300727	62.719433	29.916529	
min	0.000000	151.335266	88.050529	
25%	700.500000	245.966293	149.453659	
50%	1401.000000	279.879883	170.168365	
75%	2101.500000	330.508575	190.640427	
max	2802.000000	515.352478	258.569794	

	Thickness (depth)	Area	Perimeter	Roundness	Solidity
count	1799.000000	2803.000000	2803.000000	1946.000000	2803.000000
mean	109.705378	26511.117374	743.863770	0.470466	0.955828
std	18.940597	13782.561344	230.632076	0.118673	0.039596
min	59.494278	6037.000000	311.563489	0.173748	0.718772
25%	97.091682	16211.500000	571.730009	0.384810	0.944579
50%	110.280136	23440.500000	707.487369	0.472718	0.970422
75%	121.392773	33451.000000	878.896530	0.577553	0.981484
max	181.845200	89282.000000	1864.947387	0.697293	0.992889

	Compactness	Aspect Ratio	Eccentricity	Extent	\
count	2803.000000	1004.000000	1004.000000	2803.000000	

mean	1.825233	1.753216	0.813114	0.724587
std	0.794058	0.206616	0.041312	0.047474
min	1.164469	1.400082	0.699897	0.454538
25%	1.357398	1.612490	0.784476	0.701673
50%	1.576412	1.705716	0.810120	0.733720
75%	1.965953	1.833339	0.838141	0.757551
max	9.660057	2.731251	0.930563	0.845813

```

convex hull(convex area)
count      2803.000000
mean      27696.218159
std      14237.347610
min      6355.000000
25%     17088.500000
50%     24589.000000
75%     34863.250000
max     90642.500000
Type
SANORA      943
MAMRA      933
REGULAR     927
Name: count, dtype: int64

```

Handling missing values by filling with the mean (imputation)

Pairplot for visualizing relationships between features and 'Type'

```

data_filled = df.copy()
columns_with_missing_values = ['Length (major axis)', 'Width (minor axis)',
                               'Thickness (depth)', 'Roundness', 'Aspect Ratio', 'Eccentricity']

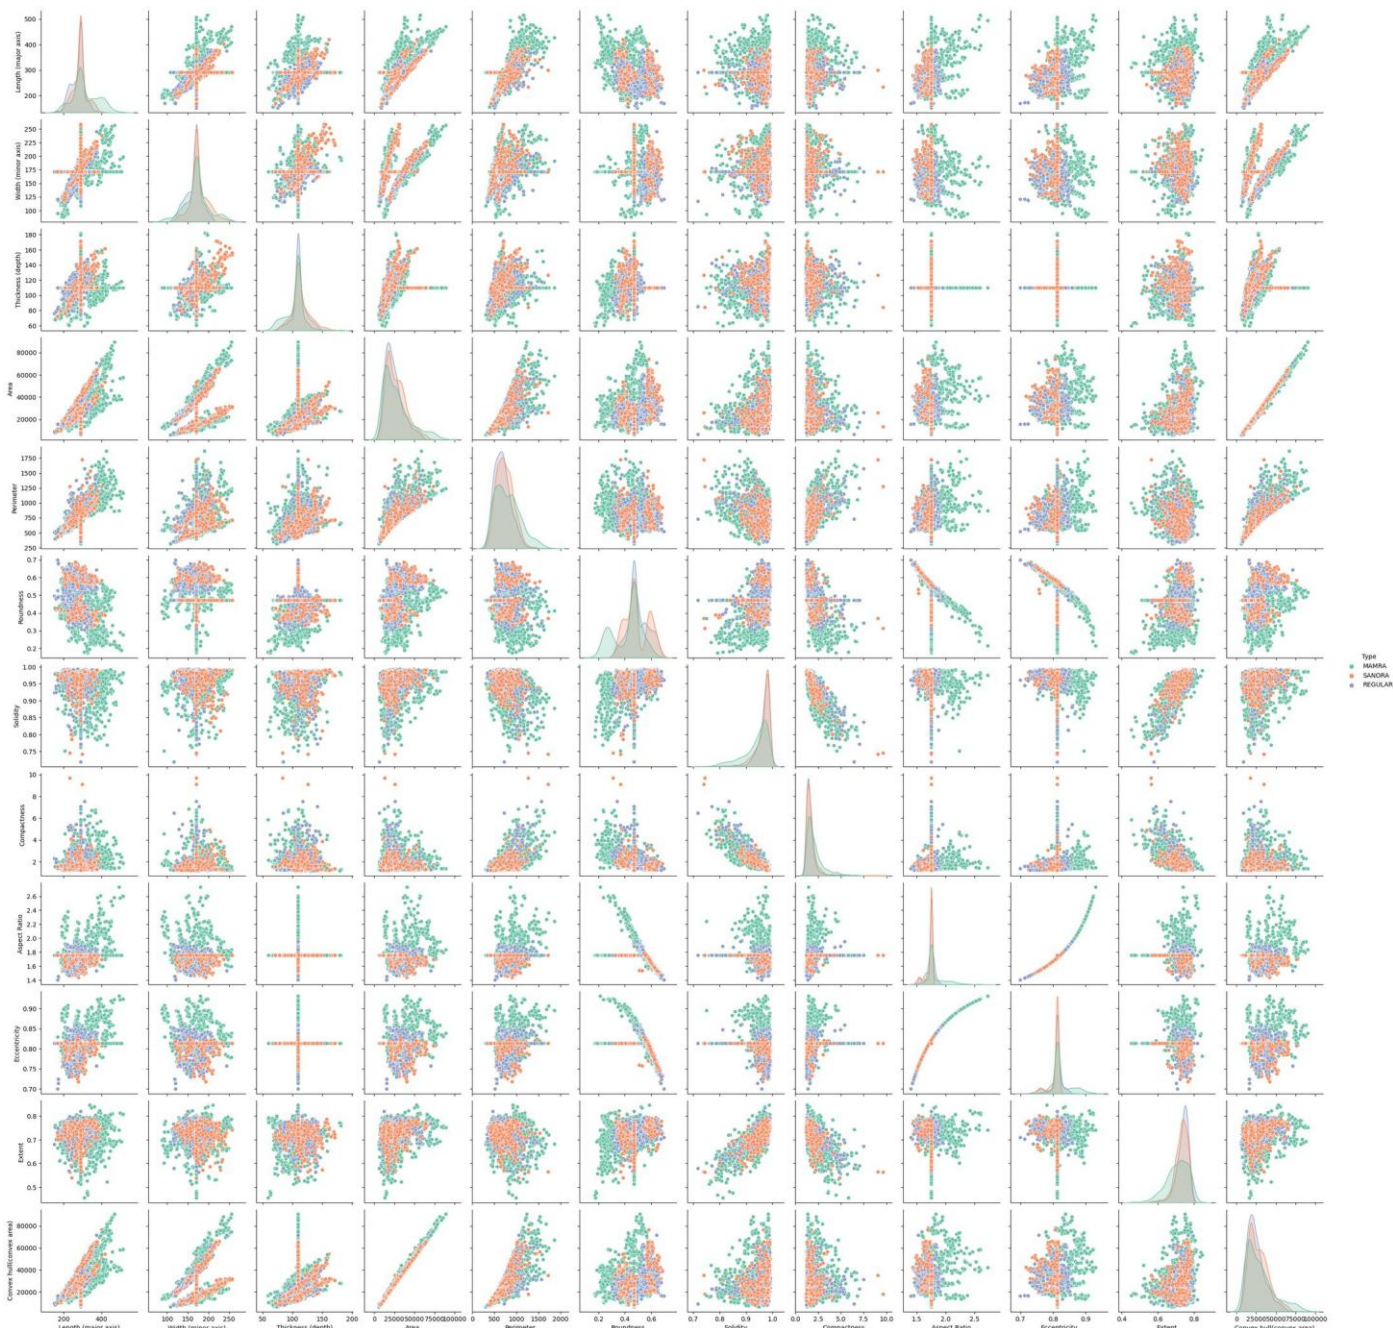
for col in columns_with_missing_values:
    data_filled[col].fillna(data_filled[col].mean(), inplace=True)

data_cleaned = data_filled.drop(columns=['Unnamed: 0'])

sns.pairplot(data_cleaned, hue='Type', diag_kind='kde', palette='Set2')
plt.show()

d:\Anaconda\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure
layout has changed to tight
    self._figure.tight_layout(*args, **kwargs)

```



Displaying head of the cleaned dataset

```
data_cleaned.head()
```

	Length (major axis)	Width (minor axis)	Thickness (depth)	Area	\		
0	290.609274	227.940628	127.759132	22619.0			
1	290.609274	234.188126	128.199509	23038.0			
2	290.609274	229.418610	125.796547	22386.5			
3	290.609274	232.763153	125.918808	22578.5			
4	290.609274	230.150742	107.253448	19068.0			
	Perimeter	Roundness	Solidity	Compactness	Aspect Ratio	Eccentricity	\
0	643.813269	0.470466	0.973384	1.458265	1.753216	0.813114	
1	680.984841	0.470466	0.957304	1.601844	1.753216	0.813114	
2	646.943212	0.470466	0.967270	1.487772	1.753216	0.813114	
3	661.227483	0.470466	0.965512	1.540979	1.753216	0.813114	
4	624.842706	0.470466	0.951450	1.629395	1.753216	0.813114	
	Extent	Convex hull (convex area)		Type			
0	0.681193	23237.5		MAMRA			

1	0.656353	24065.5	MAMRA
2	0.683620	23144.0	MAMRA
3	0.685360	23385.0	MAMRA
4	0.714800	20041.0	MAMRA

Displaying info of the cleaned dataset

```
data_cleaned.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2803 entries, 0 to 2802
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Length (major axis)                  2803 non-null   float64
1   Width (minor axis)                   2803 non-null   float64
2   Thickness (depth)                    2803 non-null   float64
3   Area                                 2803 non-null   float64
4   Perimeter                           2803 non-null   float64
5   Roundness                           2803 non-null   float64
6   Solidity                            2803 non-null   float64
7   Compactness                         2803 non-null   float64
8   Aspect Ratio                        2803 non-null   float64
9   Eccentricity                        2803 non-null   float64
10  Extent                              2803 non-null   float64
11  Convex hull (convex area)            2803 non-null   float64
12  Type                                 2803 non-null   object
dtypes: float64(12), object(1)
memory usage: 284.8+ KB
```

Encoding Categorical data.

```
le = LabelEncoder()
for column in data_cleaned.columns:
    data_cleaned[column] = le.fit_transform(data_cleaned[column])
```

Displaying Head after converting

```
data_cleaned.head(20)
```

	Length (major axis)	Width (minor axis)	Thickness (depth)	Area \
0	1138	1780	1525	1316
1	1138	1808	1539	1334
2	1138	1788	1467	1306
3	1138	1800	1473	1315
4	1138	1792	781	1020
5	1138	1796	801	1052
6	1138	1770	752	965
7	1138	1769	594	902
8	1852	947	1684	2384
9	1865	947	1566	2360
10	1885	947	1411	2315
11	1860	947	1601	2389
12	1718	947	534	1771
13	1673	947	473	1680
14	1651	947	496	1644
15	1663	947	417	1621
16	1364	903	873	2280

17	1384	886	873	2288
18	1399	883	873	2306
19	1375	870	873	2299

	Extent	Convex hull(convex area)	Type
0	420	1292	0
1	234	1335	0
2	443	1286	0
3	468	1303	0
4	931	1005	0
5	1216	1023	0
6	1215	932	0
7	1586	846	0
8	2374	2384	0
9	2135	2362	0
10	2474	2331	0
11	1822	2387	0
12	2602	1761	0
13	2099	1669	0
14	2695	1648	0
15	2541	1605	0
16	2600	2226	0
17	2529	2245	0
18	2689	2267	0
19	2494	2263	0

Displaying tail after converting

```
data_cleaned.tail(20)
```

	Length (major axis)	Width (minor axis)	Thickness (depth)	Area	\
2783	1289	1230	873	2424	
2784	1242	1047	873	2367	
2785	1234	1080	873	2379	
2786	1107	947	1374	1616	
2787	1241	947	1311	1645	
2788	998	947	1285	1515	

2789	1043	947	1219	1516
2790	1291	947	904	1573
2791	1254	947	1120	1624
2792	1212	947	923	1536
2793	1189	947	914	1513
2794	1138	1715	1501	1214
2795	1138	1661	1513	1184
2796	1138	1664	1454	1135
2797	1138	1662	1476	1151
2798	1138	1439	1383	948
2799	1138	1298	1257	799
2800	1138	1296	1271	844
2801	1138	1351	1323	885
2802	805	1066	873	2206

\	Perimeter	Roundness	Solidity	Compactness	Aspect Ratio	Eccentricity
2783	2165	1610	1193	1138	381	381
2784	2199	1480	1228	1514	515	515
2785	2127	1554	1470	1224	460	460
2786	2164	647	733	2435	622	541
2787	2028	503	523	2308	622	541
2788	1968	616	731	2344	622	541
2789	2043	573	462	2402	622	541
2790	1874	345	867	2193	622	541
2791	1892	469	711	2149	622	541
2792	1677	427	962	1913	622	541
2793	1666	435	942	1926	622	541
2794	1274	961	710	1852	622	541
2795	1218	961	896	1801	622	541
2796	1393	961	717	2092	622	541
2797	1122	961	1005	1690	622	541
2798	1110	961	483	1939	622	541
2799	745	961	861	1364	622	541
2800	887	961	789	1718	622	541
2801	999	961	701	1843	622	541
2802	2132	1872	761	1761	75	75

	Extent	Convex hull (convex area)	Type
2783	2127	2423	2
2784	1626	2358	2


2785	2197			2359	2
2786	2023			1650	2
2787	913			1699	2
2788	2166			1524	2
2789	522			1566	2
2790	741			1572	2
2791	168			1658	2
2792	532			1529	2
2793	428			1500	2
2794	799			1231	2
2795	1063			1188	2
2796	1007			1148	2
2797	1334			1136	2
2798	1161			980	2
2799	912			786	2
2800	999			827	2
2801	1524			881	2
2802	1078			2207	2

Finding correlation

```
correlation_matrix = data_cleaned.corr()
```

Plotting correlation

```
plt.figure(figsize=(25, 15))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", linewidths=0.5)
plt.title("Correlation Matrix")
plt.show()
```



Defining X ('Feature') and y ('Target')

```
X = data_cleaned.drop(columns=['Type'])
y = data_cleaned['Type']
```

Splitting Dataset into train and test

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)
```

Function to plot the confusion matrix

```
def plot_confusion_matrix(y_test, y_pred, model_name):
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=True,
yticklabels=True)
    plt.title(f"Confusion Matrix - {model_name}")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()
```

Function to plot feature importance


```
def plot_feature_importance(model, X_train, model_name):
    feature_importance = model.feature_importances_
    sorted_idx = feature_importance.argsort()
    plt.barh(X_train.columns[sorted_idx], feature_importance[sorted_idx])
    plt.xlabel("Feature Importance")
    plt.title(f"Feature Importance - {model_name}")
    plt.show()
```

Bagging Classifier

```
bagging_model = BaggingClassifier(random_state=42)
bagging_model.fit(X_train, y_train)
y_pred_bagging = bagging_model.predict(X_test)
print("Bagging Classifier Accuracy:", accuracy_score(y_test, y_pred_bagging))
print(classification_report(y_test, y_pred_bagging))
plot_confusion_matrix(y_test, y_pred_bagging, "Bagging Classifier")
```

Bagging Classifier Accuracy: 0.821640903686088

	precision	recall	f1-score	support
0	0.89	0.87	0.88	287
1	0.78	0.79	0.78	278
2	0.79	0.80	0.80	276
accuracy			0.82	841
macro avg	0.82	0.82	0.82	841
weighted avg	0.82	0.82	0.82	841


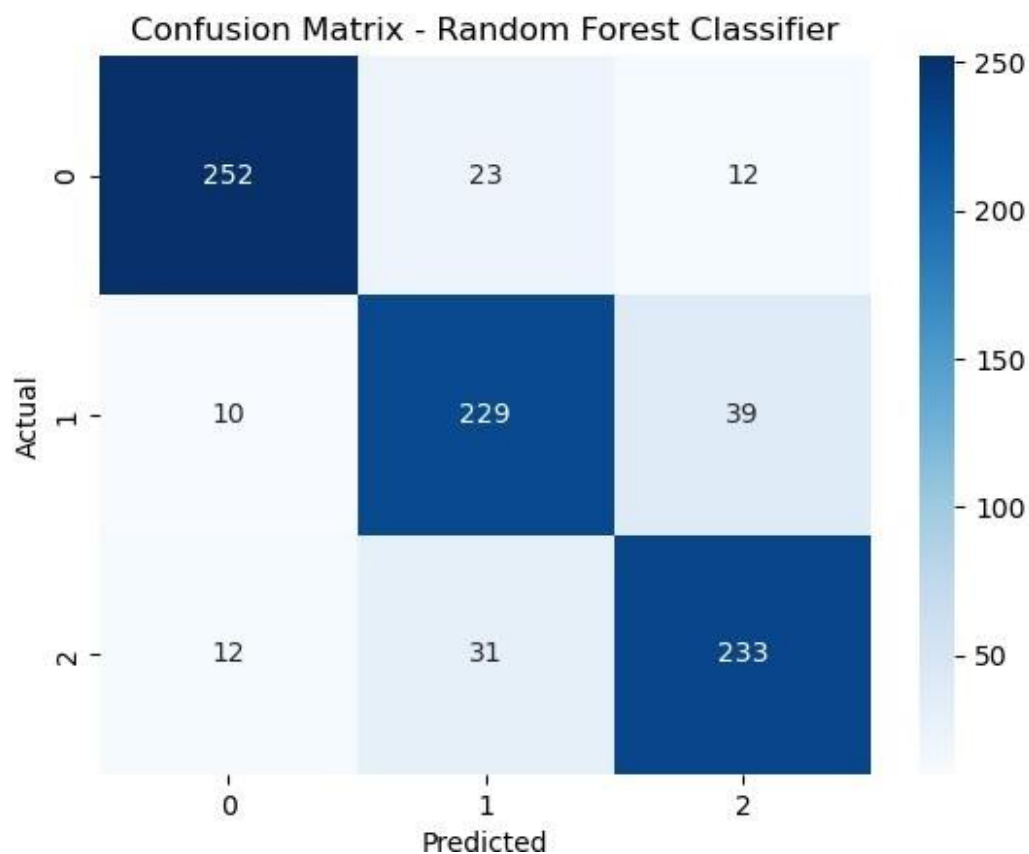


Random Forest Classifier

```
random_forest_model = RandomForestClassifier(random_state=42)
random_forest_model.fit(X_train, y_train)
y_pred_random_forest = random_forest_model.predict(X_test)
y_pred_random_forest_prob = random_forest_model.predict_proba(X_test)[:, 1]
print("Random Forest Classifier Accuracy:", accuracy_score(y_test,
y_pred_random_forest))
print(classification_report(y_test, y_pred_random_forest))
plot_confusion_matrix(y_test, y_pred_random_forest, "Random Forest
Classifier")
plot_feature_importance(random_forest_model, X_train, "Random Forest
Classifier")
```

Random Forest Classifier Accuracy: 0.8489892984542212

	precision	recall	f1-score	support
0	0.92	0.88	0.90	287
1	0.81	0.82	0.82	278
2	0.82	0.84	0.83	276
accuracy			0.85	841
macro avg	0.85	0.85	0.85	841
weighted avg	0.85	0.85	0.85	841



AdaBoost Classifier

```
adaboost_model = AdaBoostClassifier(random_state=42)
adaboost_model.fit(X_train, y_train)
y_pred_adaboost = adaboost_model.predict(X_test)
y_pred_adaboost_prob = adaboost_model.predict_proba(X_test)[: , 1]
print("AdaBoost Classifier Accuracy:", accuracy_score(y_test,
```


```

y_pred_adaboost))
print(classification_report(y_test, y_pred_adaboost))
plot_confusion_matrix(y_test, y_pred_adaboost, "AdaBoost Classifier")

```

AdaBoost Classifier Accuracy: 0.6777645659928656

	precision	recall	f1-score	support
0	0.86	0.76	0.81	287
1	0.59	0.56	0.58	278
2	0.61	0.71	0.66	276
accuracy			0.68	841
macro avg	0.69	0.68	0.68	841
weighted avg	0.69	0.68	0.68	841



Gradient Boosting Classifier

```

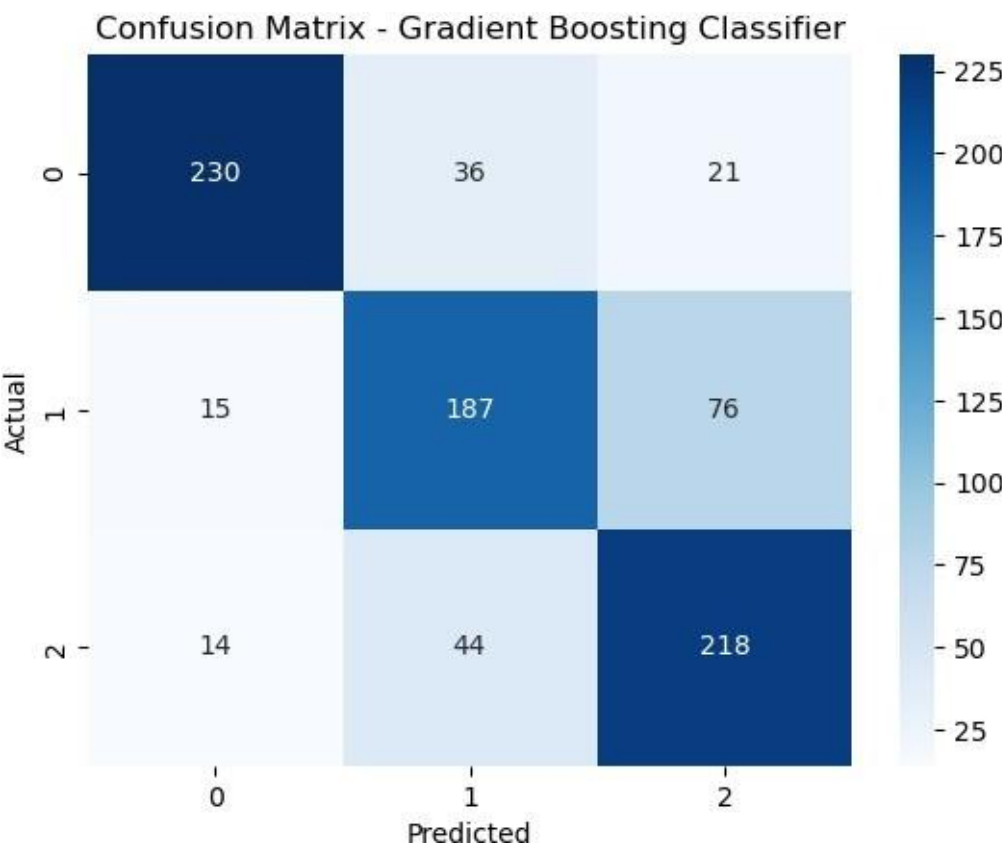
gradientboost_model = GradientBoostingClassifier(random_state=42)
gradientboost_model.fit(X_train, y_train)
y_pred_gradientboost = gradientboost_model.predict(X_test)
y_pred_gradientboost_prob = gradientboost_model.predict_proba(X_test)[:, 1]
print("Gradient Boosting Classifier Accuracy:", accuracy_score(y_test,
y_pred_gradientboost))
print(classification_report(y_test, y_pred_gradientboost))
plot_confusion_matrix(y_test, y_pred_gradientboost, "Gradient Boosting
Classifier")
plot_feature_importance(gradientboost_model, X_train, "Gradient Boosting
Classifier")

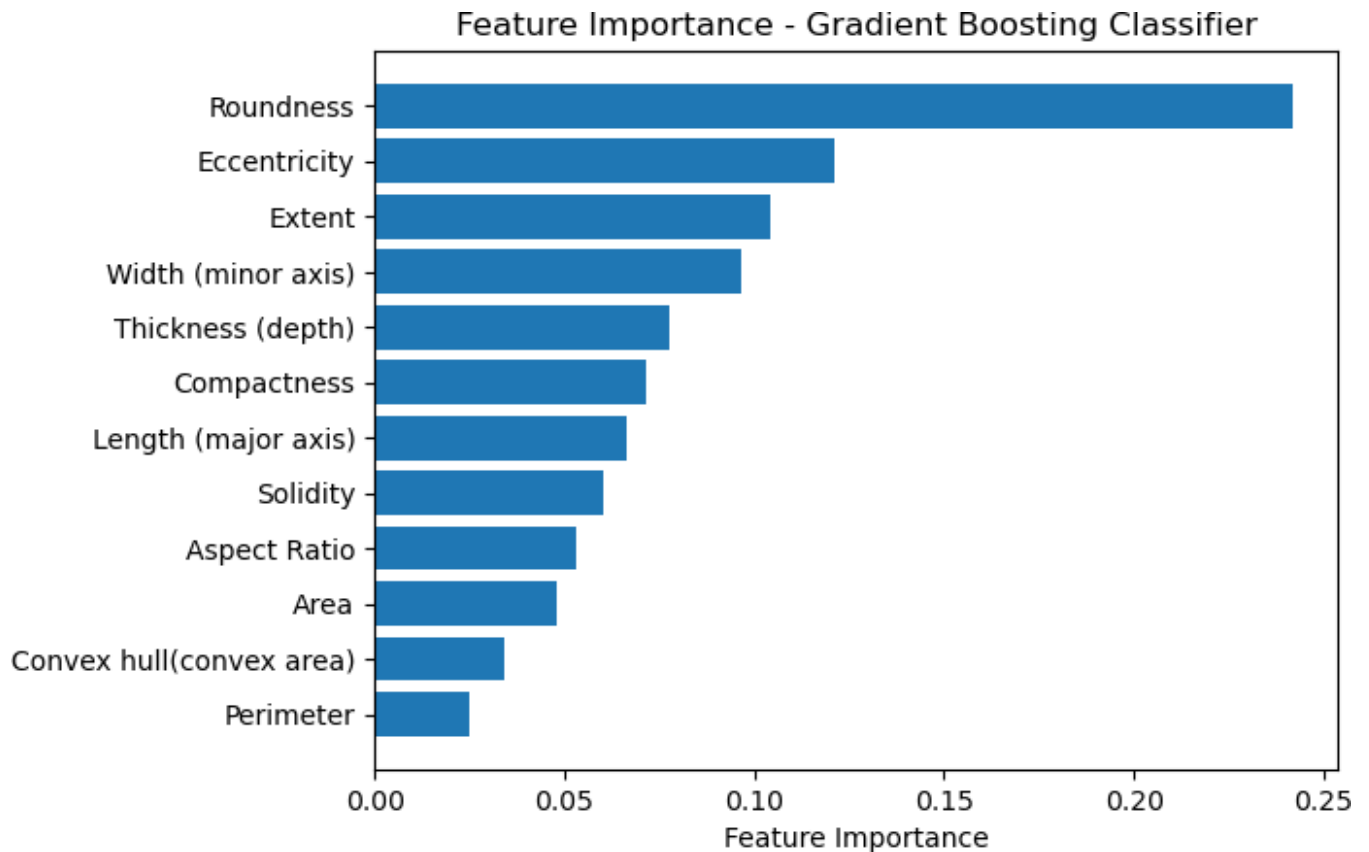
```

Gradient Boosting Classifier Accuracy: 0.7550535077288941

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.89	0.80	0.84	287
	1	0.70	0.67	0.69	278
	2	0.69	0.79	0.74	276
	accuracy			0.76	841
	macro avg	0.76	0.75	0.76	841
	weighted avg	0.76	0.76	0.76	841





Function to print summary

```
def print_model_summary(model_name, y_test, y_pred):
    print(f"Summary for {model_name}:")
    print("-----")
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
    print(f"Precision: {precision_score(y_test, y_pred,
average='weighted'):.4f}")
    print(f"Recall: {recall_score(y_test, y_pred, average='weighted'):.4f}")
    print(f"F1 Score: {f1_score(y_test, y_pred, average='weighted'):.4f}")
```

Printing summary of the model

```
# Bagging Classifier
print_model_summary("Bagging Classifier", y_test, y_pred_bagging)

# AdaBoost Classifier
print_model_summary("AdaBoost Classifier", y_test, y_pred_adaboost)

# Gradient Boosting Classifier
print_model_summary("Gradient Boosting Classifier", y_test,
y_pred_gradientboost)

# Random Forest Classifier
print_model_summary("Random Forest Classifier", y_test, y_pred_random_forest)
```

Summary for Bagging Classifier:

Accuracy: 0.8216

Precision: 0.8224

Recall: 0.8216

F1 Score: 0.8220

Summary for AdaBoost Classifier:

Accuracy: 0.6778
Precision: 0.6872
Recall: 0.6778
F1 Score: 0.6802
Summary for Gradient Boosting Classifier:

Accuracy: 0.7551
Precision: 0.7617
Recall: 0.7551
F1 Score: 0.7565
Summary for Random Forest Classifier:

Accuracy: 0.8490
Precision: 0.8506
Recall: 0.8490
F1 Score: 0.8495