

Q. Create a Multi variable Regression model of your choice using suitable dataset.

Regression Models:

1. Multiple Regression
2. Naive Bayes Regression Model
3. Naive Bayes Classification Model
4. Random Forest Classification Model

Perform EDA on the dataset and implement the model and calculate accuracy score, Precision, Recall, F1 score of each linear model

Importing all necessary library

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.tree import plot_tree
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, precision_score, recall_score,
f1_score
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, auc
```

Load the dataset and displaying head of the dataset.

0	p	x	s	n	t	p	f
1	e	x	s	y	t	a	f
2	e	b	s	w	t	l	f
3	p	x	y	w	t	p	f
4	e	x	s	g	f	n	f

	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	\
0	c	n	k	...	s	
1	c	b	k	...	s	

```

2          c          b          n  ...          s
3          c          n          n  ...          s
4          w          b          k  ...          s

stalk-color-above-ring stalk-color-below-ring veil-type veil-color \
0          w          w          p          w
1          w          w          p          w
2          w          w          p          w
3          w          w          p          w
4          w          w          p          w

ring-number ring-type spore-print-color population habitat
0          o          p          k          s          u
1          o          p          n          n          g
2          o          p          n          n          m
3          o          p          k          s          u
4          o          e          n          a          g

```

```
[5 rows x 23 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8124 entries, 0 to 8123
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null	Count	Dtype
0	class	8124	non-null	object
1	cap-shape	8124	non-null	object
2	cap-surface	8124	non-null	object
3	cap-color	8124	non-null	object
4	bruises	8124	non-null	object
5	odor	8124	non-null	object
6	gill-attachment	8124	non-null	object
7	gill-spacing	8124	non-null	object
8	gill-size	8124	non-null	object
9	gill-color	8124	non-null	object
10	stalk-shape	8124	non-null	object
11	stalk-root	8124	non-null	object
12	stalk-surface-above-ring	8124	non-null	object
13	stalk-surface-below-ring	8124	non-null	object
14	stalk-color-above-ring	8124	non-null	object
15	stalk-color-below-ring	8124	non-null	object
16	veil-type	8124	non-null	object
17	veil-color	8124	non-null	object
18	ring-number	8124	non-null	object
19	ring-type	8124	non-null	object
20	spore-print-color	8124	non-null	object
21	population	8124	non-null	object
22	habitat	8124	non-null	object

```
dtypes: object(23)
```

```
memory usage: 1.4+ MB
```

Checking null values


```
df.isnull().sum()
```

```
class          0
cap-shape      0
```

```
cap-surface      0
cap-color        0
bruises          0
odor             0
gill-attachment  0
gill-spacing     0
gill-size        0
gill-color       0
stalk-shape      0
stalk-root       0
stalk-surface-above-ring  0
stalk-surface-below-ring  0
stalk-color-above-ring   0
stalk-color-below-ring   0
veil-type        0
veil-color       0
ring-number      0
ring-type        0
spore-print-color 0
population       0
habitat          0
dtype: int64
```

Plot the distribution of the target variable 'class'

```
sns.countplot(x='class', data=df)
plt.show()
```



Encoding dataset and print head

```
le = LabelEncoder()
for column in df.columns:
    df[column] = le.fit_transform(df[column])

df.head()
```

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	\
0	1	5	2	4	1	6		1
1	0	5	2	9	1	0		1
2	0	0	2	8	1	3		1
3	1	5	3	8	1	6		1
4	0	5	2	3	0	5		1

	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	\
0	0	1	4	...	2	
1	0	0	4	...	2	
2	0	0	5	...	2	
3	0	1	5	...	2	
4	1	0	4	...	2	

	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	\
0		7	7	0	2
1		7	7	0	2
2		7	7	0	2
3		7	7	0	2
4		7	7	0	2

	ring-number	ring-type	spore-print-color	population	habitat
0	1	4	2	3	5
1	1	4	3	2	1
2	1	4	3	2	3
3	1	4	2	3	5
4	1	0	3	0	1

[5 rows x 23 columns]

Getting correlation having strongest value

```
correlation_matrix = df.corr()

# Get absolute values of correlation matrix
abs_corr_matrix = correlation_matrix.abs()

# Get pairs of variables with high correlation
strong_corrs = abs_corr_matrix.stack().reset_index()
strong_corrs.columns = ['Feature1', 'Feature2', 'Correlation']
strong_corrs = strong_corrs[strong_corrs['Feature1'] !=
strong_corrs['Feature2']]
strong_corrs = strong_corrs.sort_values(by='Correlation', ascending=False)

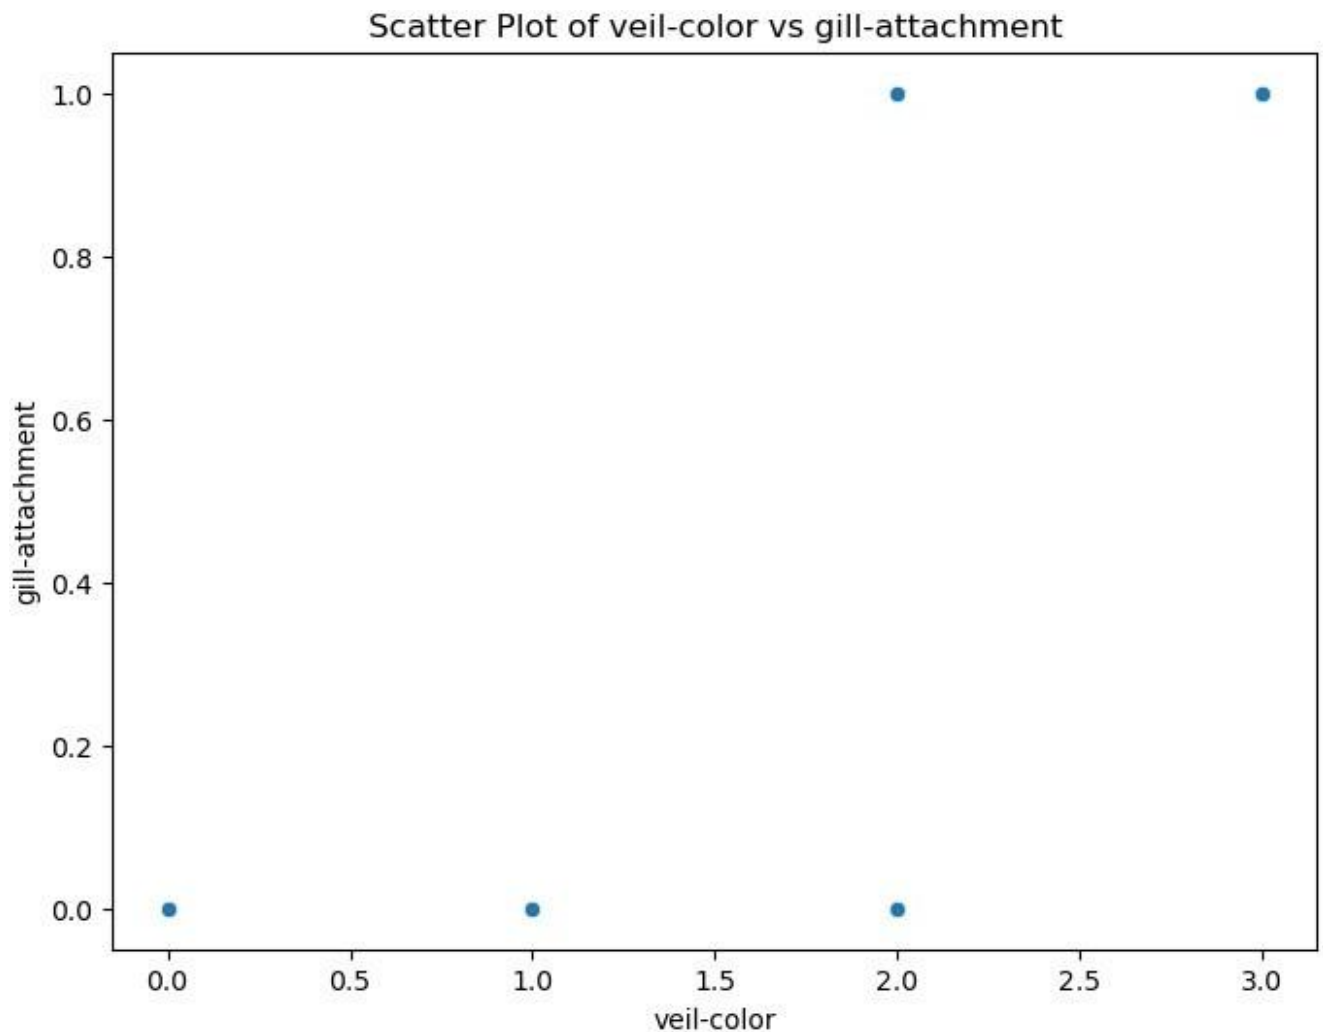
top_cor = strong_corrs.head(30)
print(top_cor)
```

	Feature1	Feature2	Correlation
358	veil-color	gill-attachment	0.897518
148	gill-attachment	veil-color	0.897518
400	ring-type	bruises	0.692973
106	bruises	ring-type	0.692973
405	ring-type	gill-color	0.629398
216	gill-color	ring-type	0.629398

Correlation having highest value

```
top_corr_pair = strong_corrs.iloc[0]
feature1 = top_corr_pair['Feature1']
feature2 = top_corr_pair['Feature2']

plt.figure(figsize=(8, 6))
sns.scatterplot(x=df[feature1], y=df[feature2])
plt.title(f'Scatter Plot of {feature1} vs {feature2}')
plt.xlabel(feature1)
plt.ylabel(feature2)
plt.show()
```



Feature and target for regression model

```
features = ['gill-spacing', 'ring-type', 'gill-size', 'gill-color']
target = 'class'

X = df[features]
y = df[target]

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

Multiple Regression


```
# Multiple Regression
model_reg = LinearRegression()
model_reg.fit(X_train, y_train)
y_pred_reg = model_reg.predict(X_test)

print("Multiple Regression Evaluation:")
print(f"R-squared: {r2_score(y_test, y_pred_reg):.2f}")
print(f"MAE: {mean_absolute_error(y_test, y_pred_reg):.2f}")
print(f"MSE: {mean_squared_error(y_test, y_pred_reg):.2f}")
print(f"RMSE: {mean_squared_error(y_test, y_pred_reg, squared=False):.2f}")
```

```
Multiple Regression Evaluation:
R-squared: 0.48
MAE: 0.30
MSE: 0.13
RMSE: 0.36
```

Scatter Plot of Actual vs. Predicted Values

```
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_reg, edgecolors=(0, 0, 0))
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--',
lw=4)
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs Predicted Values')
plt.show()
```




Residual Plot

```

residuals = y_test - y_pred_reg
plt.figure(figsize=(10, 6))
plt.scatter(y_pred_reg, residuals, edgecolors=(0, 0, 0))
plt.hlines(y=0, xmin=y_pred_reg.min(), xmax=y_pred_reg.max(), colors='r',
linestyles='dashed')
plt.xlabel('Predicted')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()

```

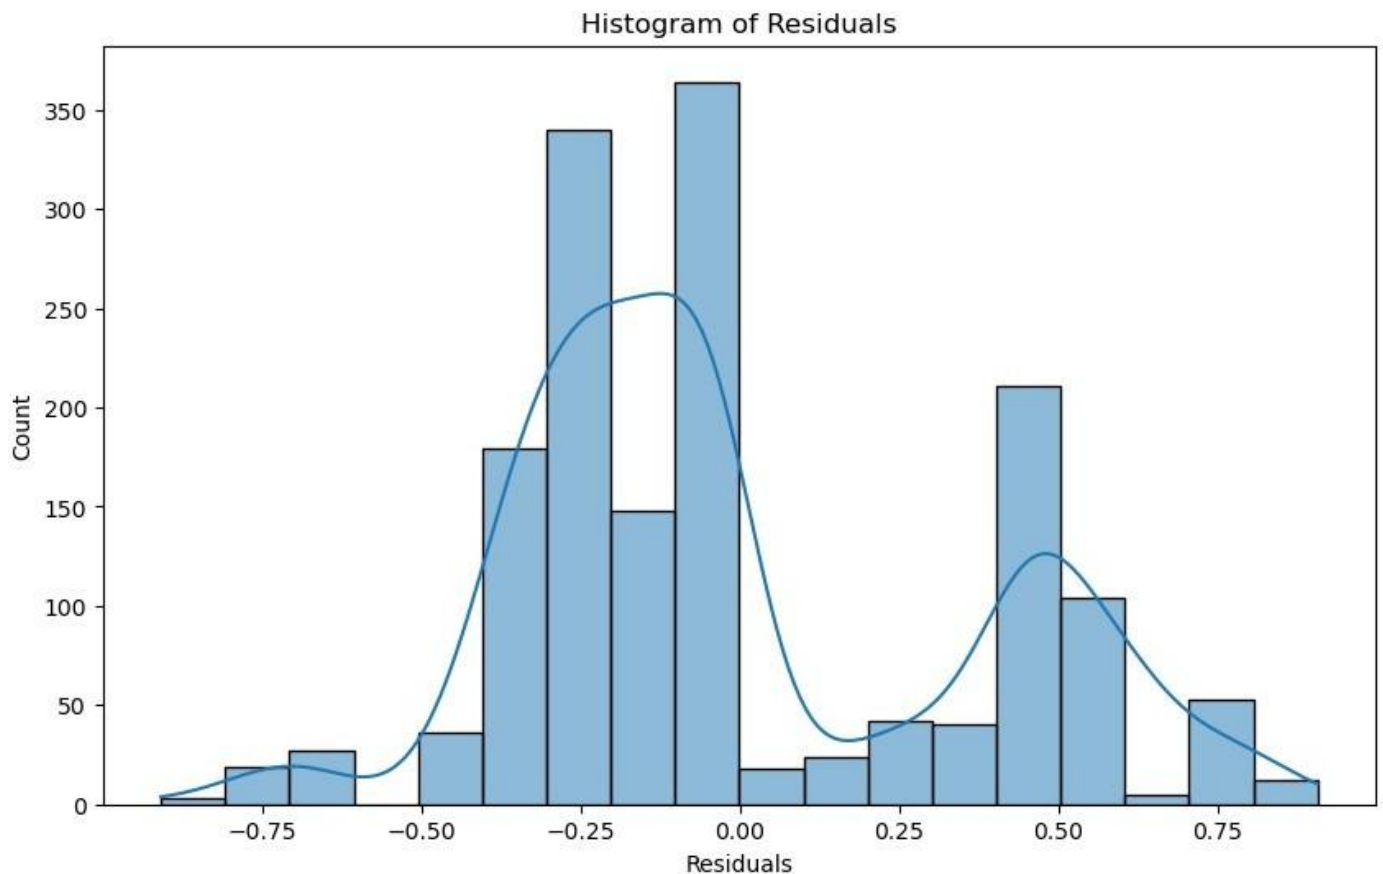


Histogram of Residuals

```

plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.xlabel('Residuals')
plt.title('Histogram of Residuals')
plt.show()

```

Decision-Tree Regression

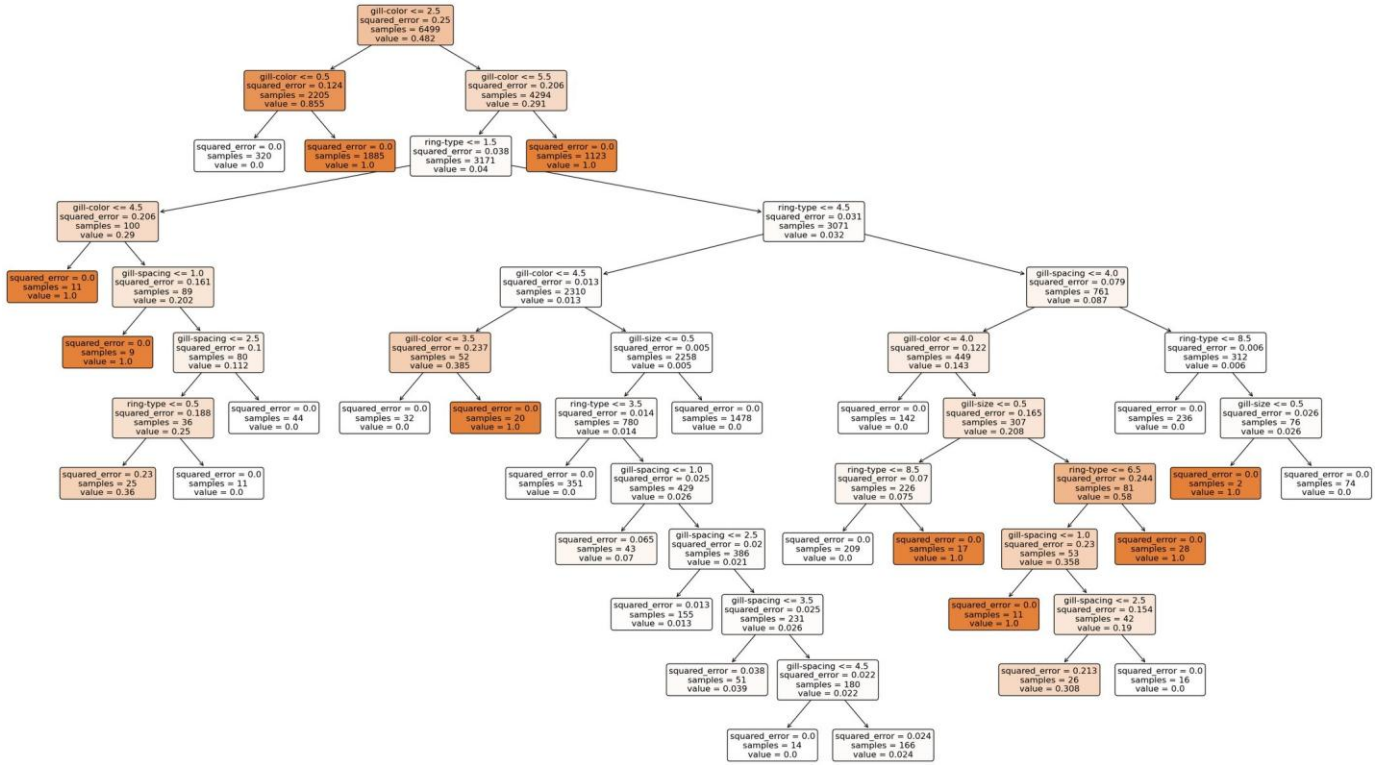
```
model_tree = DecisionTreeRegressor(random_state=42)
model_tree.fit(X_train, y_train)
y_pred_tree = model_tree.predict(X_test)

# Evaluation
print("Decision Tree Regression Evaluation:")
print(f"R-squared: {r2_score(y_test, y_pred_tree):.2f}")
print(f"MAE: {mean_absolute_error(y_test, y_pred_tree):.2f}")
print(f"MSE: {mean_squared_error(y_test, y_pred_tree):.2f}")
print(f"RMSE: {mean_squared_error(y_test, y_pred_tree, squared=False):.2f}")
```

```
Decision Tree Regression Evaluation:
R-squared: 0.98
MAE: 0.01
MSE: 0.00
RMSE: 0.07
```

The decision tree plot


```
# Plotting the decision tree
plt.figure(figsize=(35, 20))
plot_tree(model_tree, feature_names=features, filled=True, rounded=True,
          fontsize=12)
plt.title('Decision Tree Regression Model')
plt.show()
```



Feature Importances Plot

```
importances = model_tree.feature_importances_
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(10, 6))
plt.title('Feature Importances')
plt.bar(range(X_train.shape[1]), importances[indices], align='center')
plt.xticks(range(X_train.shape[1]), [features[i] for i in indices],
rotation=90)
plt.xlim([-1, X_train.shape[1]])
plt.show()
```



Feature and target for classification model

```
classification_features = ['cap-shape', 'cap-color', 'bruises', 'odor']
classification_target = 'class'

classification_X = df[classification_features]
classification_y = df[classification_target]

X_train, X_test, y_train, y_test = train_test_split(classification_X,
                                                    classification_y, test_size=0.2, random_state=42)
```

Naive Bayes Classification

```
# Naive Bayes Classification
model_nb = GaussianNB()
model_nb.fit(X_train, y_train)
y_pred_nb = model_nb.predict(X_test)

print("\nNaive Bayes Classification Evaluation:")
print(f"Accuracy Score: {accuracy_score(y_test, y_pred_nb):.2f}")
print(f"Precision: {precision_score(y_test, y_pred_nb,
                                   average='weighted'):.2f}")
print(f"Recall: {recall_score(y_test, y_pred_nb, average='weighted'):.2f}")
print(f"F1 Score: {f1_score(y_test, y_pred_nb, average='weighted'):.2f}")
```


```
Naive Bayes Classification Evaluation:
Accuracy Score: 0.73
Precision: 0.74
```

Recall: 0.73
F1 Score: 0.73

Confusion Matrix for Naive Bayes classification model.

```
conf_matrix = confusion_matrix(y_test, y_pred_nb) # Replace y_pred_rf with
your predictions

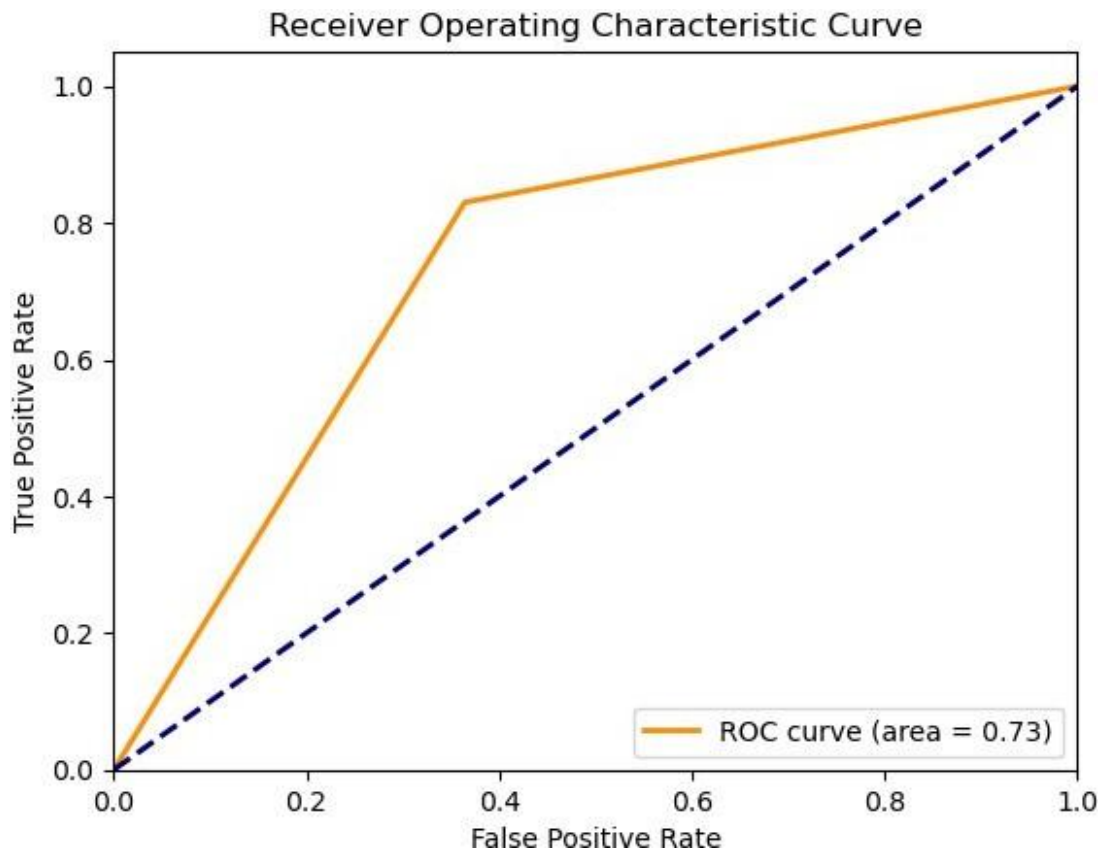
plt.figure()
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



The ROC curve (Receiver Operating Characteristic curve) for Naive Bayes classification model.

```
fpr, tpr, _ = roc_curve(y_test, y_pred_nb)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic Curve')
plt.legend(loc='lower right')
plt.show()
```



Random Forest Classification

```
# Random Forest Classification
model_rf = RandomForestClassifier(n_estimators=100, random_state=42)
model_rf.fit(X_train, y_train)
y_pred_rf = model_rf.predict(X_test)


print("\nRandom Forest Classification Evaluation:")
print(f"Accuracy Score: {accuracy_score(y_test, y_pred_rf):.2f}")
print(f"Precision: {precision_score(y_test, y_pred_rf,
average='weighted'):.2f}")
print(f"Recall: {recall_score(y_test, y_pred_rf, average='weighted'):.2f}")
print(f"F1 Score: {f1_score(y_test, y_pred_rf, average='weighted'):.2f}")
```

```
Random Forest Classification Evaluation:
Accuracy Score:0.99
Precision: 0.99
Recall: 0.99
F1 Score: 0.99
```

Confusion Matrix for RandomForest classification model.

```
conf_matrix = confusion_matrix(y_test, y_pred_rf)

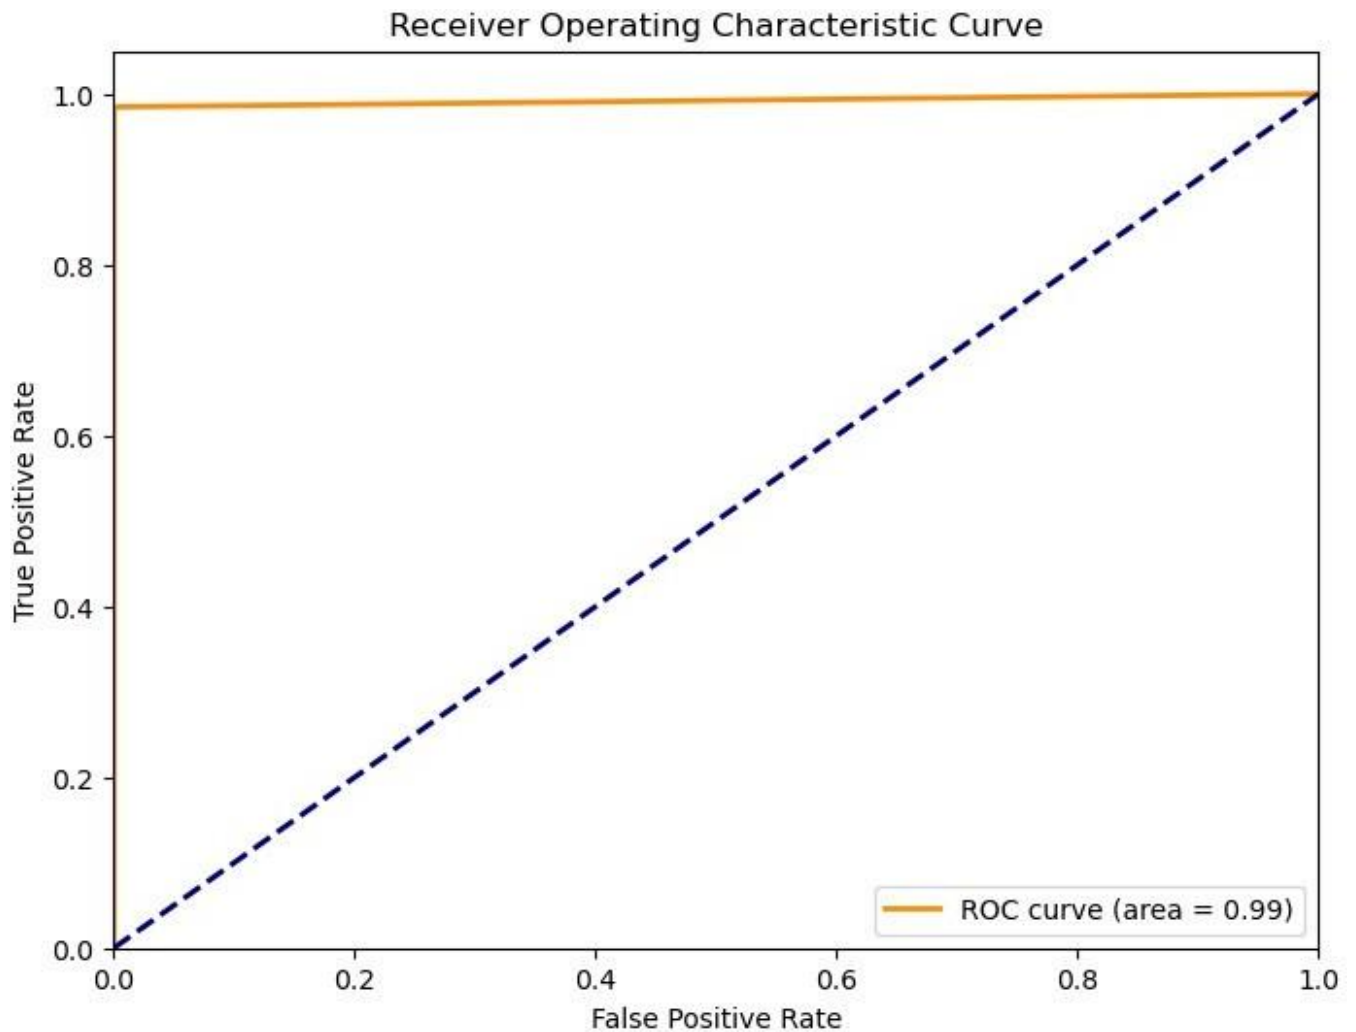
plt.figure()
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



The ROC curve (Receiver Operating Characteristic curve) for RandomForest classification model.

```
fpr, tpr, _ = roc_curve(y_test, y_pred_rf)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic Curve')
plt.legend(loc='lower right')
plt.show()
```



Feature Importances Plot (For tree-based classifiers like Random Forest, you can plot the importance of features.)

```
importances = model_rf.feature_importances_  
indices = np.argsort(importances)[::-1]  
  
plt.figure(figsize=(10, 6))  
plt.title('Feature Importances')  
plt.bar(range(X_train.shape[1]), importances[indices], align='center')  
plt.xticks(range(X_train.shape[1]), [features[i] for i in indices],  
rotation=90)  
plt.xlim([-1, X_train.shape[1]])  
plt.show()
```

Feature Importances

